

This section appears to be a collection of notes and information about a language model called Llama 2, which is a type of large language model (LLM) developed by Meta AI. Here's a summary of the key points:

Llama 2 Overview

- Llama 2 is an auto-regressive language model that uses an optimized transformer architecture.
- It was trained on 2 trillion tokens of data from publicly available sources, with fine-tuning on over one million new human-annotated examples.

Model Details

- Llama 2 has three different parameter sizes: 7B, 13B, and 70B.
- The model can be used for both text input and output.
- It is designed to be helpful and safe, with a focus on assisting users with tasks such as chat and language generation.

Training Factors

- The model was trained using custom training libraries, Meta's Research Super Cluster, and production clusters for pretraining.
- Fine-tuning, annotation, and evaluation were also performed on third-party cloud compute.
- The estimated carbon footprint of the model is 539 tCO₂eq, which has been offset by Meta's sustainability program.

Ethical Considerations

- Llama 2 carries risks with use, and testing conducted to date has only covered English scenarios.
- Developers should perform safety testing and tuning tailored to their specific applications of the model before deploying it.

Intended Use Cases

- Llama 2 is intended for commercial and research use in English.
- Tuned models are intended for assistant-like chat, while pretrained models can be adapted for a variety of natural language generation tasks.

Use Restrictions

- The model should not be used in any manner that violates applicable laws or regulations, including trade compliance laws.
- Use in languages other than English is prohibited.
- Other uses may be prohibited by the Acceptable Use Policy and Licensing Agreement for Llama 2.